



❖ General goal


- ❖ Advance the state-of-the-art in **predictive simulation**.
- ❖ Connect **quantum / molecular simulations** of materials with state-of-the-art **programming languages, compilers, and hpc**.
- ❖ **Uncertainty quantification**

❖ Motivating problem

- ❖ Predict the degradation of complex materials under extreme loading, inaccessible to direct experimental observation
- ❖ In particular: **ultrahigh temperature ceramics in hypersonic flows**.


❖ Supported by: Department of Energy's Predictive Science Academic Alliance Program (PSAAP-III)

❖ Interdisciplinary project: CCSE, CSAIL, SDSC, external collaborations




❖ General goals

- ❖ ***Facilitate the development** of chemical and materials simulations in **CESMIX***
- ❖ *Demonstrate that it is possible to address these simulations using  
**a single high-level programming language without compromising performance***
- ❖ ***Expand** the capabilities of the **atomistic ecosystem in Julia**, as well as its user **community***
- ❖ *Contribute to a **better integration** of the atomistic software ecosystem*



❖ Specific goals

- ❖ *Composable workflows* and new software *abstractions*
- ❖ *UQ-driven integration*
- ❖ *AD enabling UQ, training, and force calculations*
- ❖ Simplify the *definition, fitting, and integration of new interatomic potentials*
- ❖ Integrate with state-of-the-art tools. E.g. **LAMMPS**.
- ❖ Leverage **Julia**!



# Software roadmap

## Year 1

## Year 2

## Year 3

## Years 4–5

### CS tools

OpenCilk  
Enzyme  
Julia + OpenCilk  
Julia + FluxRM  
Julia + Spindle  
Tiramisu

Initial stable releases  
Forward mode/mixed mode  
Start integration effort  
Integration (with LLNL)  
Integration  
Initial Exasim application

Target new runtimes and backends  
Parallel-specific, Julia parallelism  
Compiler and runtime integration  
Initial design for UQ workflows

New tools using CSI  
BLAS and split forward mode  
Performance eng. of integration

New parallel constructs in Tapir  
Checkpointing

Adapt to DFT use cases

### Interfaces/ abstractions

Julia + HPC arch  
LAMMPS.jl  
Atomistic.jl/etc

Tapir IR extension, HPC readiness  
Initial design of LAMMPS.jl  
Initial design

Parallel optimization, parallel runtimes  
Enzyme for AD in LAMMPS  
Complete coverage; interface  
with ACE.jl, ASE, KMC.jl

Perf opt for specific HPC archs  
FluxRM.jl support for key  
workflows

### DFT

DFTK.jl

Forward-mode AD

Reverse AD; sensitivity analysis.  
DFTK.jl on GPUs

Expose multi-fidelity methods;  
GPU performance optimization

Explore DSL (Tiramisu)  
opportunities

### MD + potentials

MDP  
Julia atomistic suite  
KMC

Initial design & development  
Initial design & development  
Initial design & development

Julia interfaces; profiling + optimization  
Integration of all potentials, DFT  
codes, MD simulators, KMC

Automated UQ + active learning  
workflows for MD; integration  
with FluxRM.jl

Integration of MD UQ workflow  
with KMC and fluids simulations

### fluids

Exasim

Exasim + Tiramisu; Julia interface

Exasim + Enzyme;  
Integration with Mutation++

Performance optimization:  
chemistry on GPUs

Surface chemistry/erosion  
models

### UQ

CESMIX-UQ

Bayesian potential fitting  
(PotentialUQ.jl)

UQ for MD simulations  
Initial active learning workflows

UQ + multi-fidelity for DFT  
Parametric UQ for Exasim

Coupled UQ workflows (Y4: only  
molecular; Y5: including fluids)

# Integrating atomistic software...

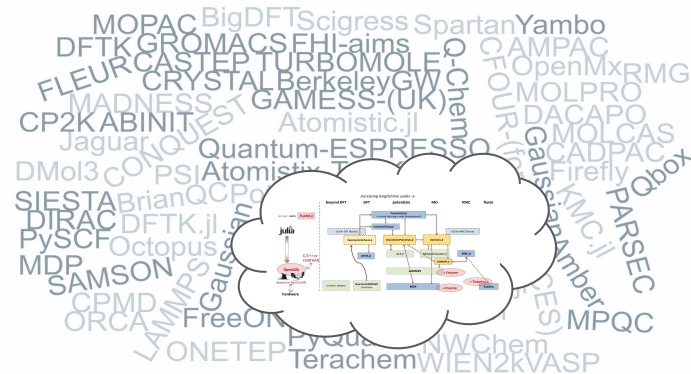
❖ Many computational tools under development dedicated to **chemistry** and **materials science**

❖ Integration is a must!

- ❖ **Broadens** the spectrum of possible workflows
- ❖ **Manage** problems arising from the **interactions of components**  
E.g. **competition for hardware resources**
- ❖ **Catalyze advances** in chemistry and materials science

❖ But it is also a **challenge**...

- ❖ Software development is becoming more and more accessible
  - Increase in the number of tools
  - Increase in complexity associated with **overlapping functionality**
- ❖ Horizontal hierarchies and large communities make it **difficult to reach consensus**



❖ Current integration efforts... ASE (Atomic Simulation Environment)



### ❖ **Fast & easy :-D**

- ❖ Addresses the **two language problem**: Prototype algorithms in a user-friendly language, and then have to rewrite them in a faster language.
- ❖ More **control** over your work
- ❖ **Abstracts performance** management as much as possible
- ❖ Avoid duplication of effort
- ❖ More time spent on your research rather than on performance
- ❖ Optimized code is easy to read
- ❖ Collaboration and integration are easier

### ❖ **Software composability** through multiple dispatch

- ❖ Accelerate the development of **atomistic workflows**.
- ❖ Accelerate the development of **machine learning interatomic potentials** defined by the composition of neural networks with state-of-the-art interatomic potential descriptors.
- ❖ Allows it to express **object-oriented** and **functional** programming patterns

### ❖ **SciML**

- ❖ Open Source Software for Scientific Machine Learning.
- ❖ E.g. **Optimization.jl**: brings together several **optimization packages into one unified Julia interface**.

### ❖ **Automatic differentiation**

- ❖ Accelerate development by automatically calculating **loss function gradients, forces, sensitivity of QoI**, etc.
- ❖ Allows to calculate **gradients across multiple languages and simulation codes**.
- ❖ Powerful tools, such as **Enzyme.jl** or **Zygote.jl**.

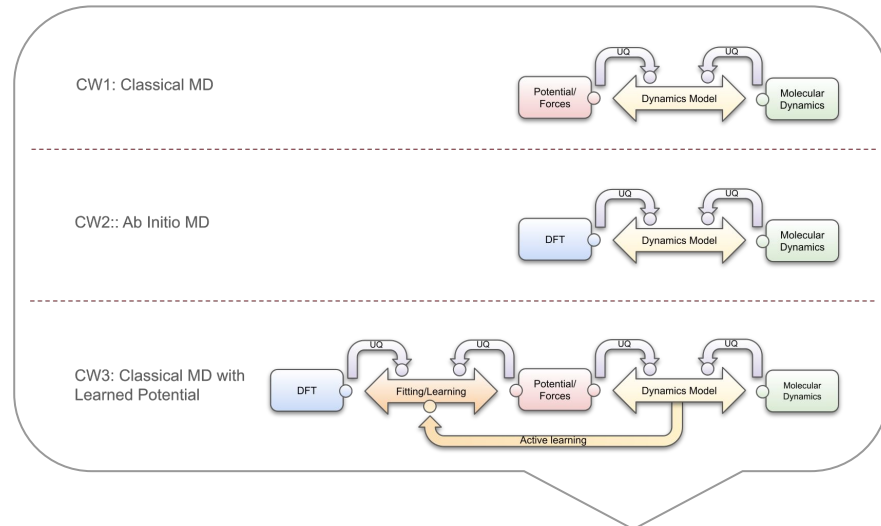
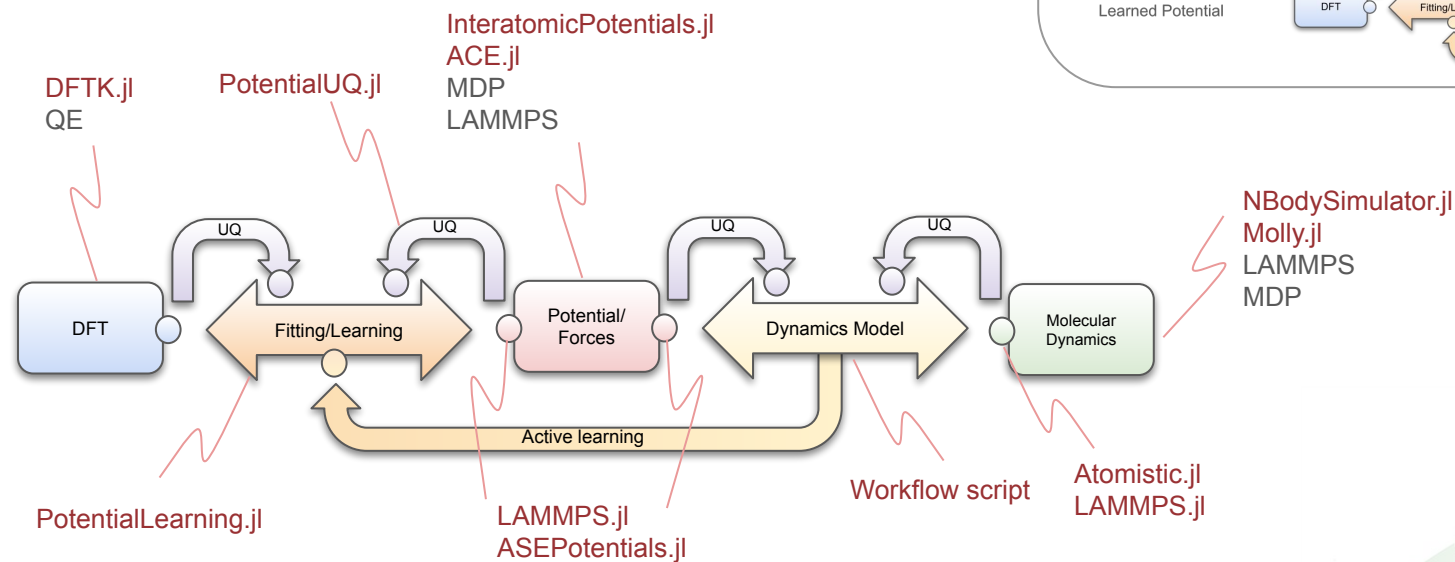
### ❖ **HPC and Machine learning** abstractions

- ❖ Accelerate the development and execution of **machine learning interatomic potentials**. **Flux.jl** makes parallel learning simple using the NVIDIA GPU abstractions of **CUDA.jl**.


### ❖ **More!**

- ❖ Dynamically typed, but with support for **optional type declarations**.
- ❖ **Feels like a scripting language, but can be compiled to efficient native code** for multiple platforms via LLVM.
- ❖ ...

## Integrating atomistic software in CESMIX: Composable Workflows



## Composable Workflows



Only if needed: ☐ ☐ ☐ ☐ ☐



❖ **AtomsBase.jl**: *abstract interface for representation of **atomic geometries in Julia**. It aims to be a lightweight means of facilitating interoperability between various tools including...*


- ❖ Chemical simulation engines. E.g. DFT, MD, etc.
- ❖ File I/O with standard formats (.cif, .xyz, ...)
- ❖ Numerical tools: sampling, integration schemes, etc.
- ❖ Automatic differentiation and ML systems
- ❖ Visualization (e.g. plot recipes)

❖ **AtomIO.jl**: *standard **IO** package for atomic structures integrating with FileIO, AtomsBase, and others.*

```
struct Atom{D, L<:Unitful.Length, V<:Unitful.Velocity, M<:Unitful.Mass}
    position::SVector{D, L}
    velocity::SVector{D, V}
    atomic_symbol::Symbol
    atomic_number::Int
    atomic_mass::M
    data::Dict{Symbol, Any} # Store arbitrary data about the atom.
end

struct FlexibleSystem{D, S, L<:Unitful.Length} <: AbstractSystem{D}
    particles::AbstractVector{S}
    box::SVector{D, SVector{D, L}}
    boundary_conditions::SVector{D, BoundaryCondition}
end

struct FastSystem{D, L<:Unitful.Length, M<:Unitful.Mass} <: AbstractSystem{D}
    box::SVector{D, SVector{D, L}}
    boundary_conditions::SVector{D, BoundaryCondition}
    positions::Vector{SVector{D, L}}
    atomic_symbols::Vector{Symbol}
    atomic_numbers::Vector{Int}
    atomic_masses::Vector{M}
end
```



# DFTK.jl: Density-functional toolkit

- **julia** code for plane-wave DFT, started in 2019
- Fully **composable** with **julia** ecosystem:
  - Algorithmic differentiation (AD)
  - Numerical error control
  - GPU acceleration (GSoC confirmed) **planned**
  - Uncertainty quantification (UQ) **planned**
- Key tool across domains and their use cases:
  - **Mathematical analysis** (reduced models)
  - Scale-up to **applications** ( $\simeq 1000$  electrons)
  - Features incl. meta-GGA, response, MPI parallelisation
  - Speed within factor 2–4 to established codes

⇒ Build to enable **multidisciplinary synergies**

- Low entrance barrier **across backgrounds**:
  - Only 7000 lines of code, open-source components
  - Avoids **two-language problem**: Just **julia**
- **High-productivity** research framework:
  - GSoC student (10 weeks) for initial AD support



high-performance computing

materials simulations




DFTK

numerical analysis

novel scientific models



$$H\Psi = E\Psi$$



# InteratomicPotentials.jl

```
using AtomsBase, Unitful, UnitfulAtomic
# Define an atomic system
element = :Ar
atom1    = Atom(element, ( @SVector [1.0, 0.0, 0.0] ) * 1u"Å")
atom2    = Atom(element, ( @SVector [1.0, 0.25, 0.0] ) * 1u"Å")
box = [[1., 0.0, 0.0], [0.0, 1.0, 0.0], [0.0, 0.0, 1.0]] * 1u"Å"
bcs = [DirichletZero(), Periodic(), Periodic()]
system  = FlexibleSystem(atoms, box , bcs)
```

```
ϵ = 1.0 * 1u"eV"
σ = 0.25 * 1u"Å"
rcutoff = 2.25 * 1u"Å"
lj       = LennardJones(ϵ, σ, rcutoff, [element])
```

```
pe      = potential_energy(system, lj)
f       = force(system, lj)
v       = virial(system, lj)
v_tensor = virial_stress(system, lj)
```

```
EmpiricalPotential <: AbstractPotential
BornMayer <: EmpiricalPotential
LennardJones <: EmpiricalPotential
Coulomb <: EmpiricalPotential
ZBL <: EmpiricalPotential
```

```
LinearCombinationPotential <: MixedPotential
```

```
# See InteratomicBasisPotentials.jl
BasisPotential <: AbstractPotential
SNAP <: BasisPotential
ACE <: BasisPotential
```

## ❖ General goals

❖ *Manipulate interatomic potentials in Julia. Allows to compute **energies, forces, and virial tensors** for a variety of interatomic potentials.*

## ❖ InteratomicPotentials.jl/InteratomicBasisPotentials.jl

❖ *Lennard Jones, Born Mayer, Coulomb, ZBL, Morse.*

❖ **SNAP: new Julia implementation!**

❖ **ACE: interface to ACE1.jl**

## ❖ Currently working on...

❖ *Adding additional potentials*

❖ **Potentials for molecules, EAM, 3-, 4-body potentials**

❖ **Composition of neural networks with IAP descriptors**

❖ **Efficient neighbor lists**


❖ **GPU parallelization**

❖ *Generic GPU kernels for potentials*

❖ **SNAP, ACE**

## ❖ Main dependencies within this project

❖ **AtomsBase.jl**



# Create your own potential...

- ❖ **Composability:** *high level software that is readable, flexible, and **extensible**!*
- ❖ **Leverage multiple dispatch:** add functionality to external software without modifying it
  - ❖ **Define new types on which existing functions can be applied**
  - ❖ **Define new functions which can be applied on existing types**

```
struct SimpleNNPotential <: AbstractPotential
    nn
    nn_params
end

function potential_energy(R::AbstractFloat,
                        p::SimpleNNPotential)

    return p.nn(R)
end

function force(R::AbstractFloat,
              r::SVector,
              p::SimpleNNPotential)

    dnndr = first(gradient((x)->p.nn(norm(x)), r))
    return -dnndr
end
```

```
model = Chain(Dense(1,32,Flux.relu),
              Dense(32,32,Flux.relu), Dense(32,1))

nn(R) = first(model([R]))
nn_params = Flux.params(model)

p = SimpleNNPotential(nn, nn_params)

potential_energy(1.0, p)

r_diff = SVector(1.0, 1.0, 1.0)
force(norm(r_diff), r_diff, p)
```

# PotentialLearning.jl

```
# Abstract types #####

"""
    LearningProblem{D, T}
"""
abstract type LearningProblem{D, T} end

"""
    LearningOptimizer{D, T}
"""
abstract type LearningOptimizer{D, T} end

# Optimizers #####

"""
    QRLinearOpt{D, T}
"""
QR optimizer
"""
struct QRLinearOpt{D, T} <: LearningOptimizer{D, T} end

# Learning and loss functions #####

"""
    loss(params::Vector{T}, opt::LearningOptimizer{D, T})

`params`: parameters to be fitted
`opt`: learning optimizer
"""
function loss(params::Vector{T}, opt::LearningOptimizer{D, T}) where {T, D} end

"""
    learn(lp::LearningProblem{D, T}, opt::LearningOptimizer{D, T})

`lp`: learning problem
`opt`: learning optimizer
"""
function learn(lp::LearningProblem{D, T}, opt::LearningOptimizer{D, T}) where {T, D} end
```

## ❖ General goals


- ❖ Implement an open source Julia library for **active learning** of interatomic potentials in atomistic simulations of materials.
- ❖ Incorporates elements of **bayesian inference**, **machine learning**, **differentiable programming**, **software composability**, and **high-performance computing**.

## ❖ Specific goals

- ❖ Intelligent **data subsampling**. Via DPP, clustering.
- ❖ Inference of the **optimal values** and **uncertainties** of the model parameters, to propagate them through the atomistic simulation.
  - ❖ Interatomic potential **hyper-parameter** optimization.
  - ❖ Interatomic **potential fitting**.  
The potentials addressed in this package are defined in *InteratomicPotentials.jl* and *InteratomicBasisPotentials.jl*.
- ❖ Measurement of QoI **sensitivity** to individual parameters.
- ❖ **Input data** management and **post-processing**.

## ❖ Main dependencies within this project

- ❖ *AtomsBase.jl*, *InteratomicPotentials.jl*, *InteratomicBasisPotentials.jl*, *Carom.jl*, *Turing.jl*, *Enzyme.jl*.



## Classical MD

*Atomistic.jl*

```
using Atomistic
using InteratomicPotentials
using NBodySimulator
```

```
N = 864
element = :Ar
box_size = 3.47786u"nm"
reference_temp = 94.4u"K"
thermostat_prob = 0.1 # this number was chosen arbitrarily
Δt = 1e-2u"ps"

initial_system = generate_atoms_in_cubic_cell(N, element, box_size, reference_temp)
```

```
eq_steps = 2000
eq_thermostat = NBodySimulator.AndersenThermostat(austrip(reference_temp),
                                                    thermostat_prob / austrip(Δt))
eq_simulator = NBSimulator(Δt, eq_steps, thermostat = eq_thermostat)
potential = InteratomicPotentials.LennardJones(austrip(1.657e-21u"J"),
                                                austrip(0.34u"nm"),
                                                austrip(0.765u"nm"), [:Ar])
```

```
eq_result = simulate(initial_system, eq_simulator, potential)
```

```
prod_steps = 5000
prod_simulator = NBSimulator(Δt, prod_steps, t₀ = get_time(eq_result))
```

```
prod_result = simulate(get_system(eq_result), prod_simulator, potential)
```

❖ **General goal:** *provide an integrated workflow for MD simulations.*

❖ **Interface to**


❖ *Molly.jl*

❖ *NBodySimulator.jl*

❖ **Main dependencies within this project**

❖ *AtomsBase.jl*, *InteratomicPotentials.jl*, and *InteratomicBasisPotentials.jl*





## Classical MD

*Atomistic.jl*

```
using Atomistic
using InteratomicPotentials
using Molly
```

```
N = 864
element = :Ar
box_size = 3.47786u"nm"
reference_temp = 94.4u"K"
coupling_factor = 10 # this number was chosen arbitrarily
Δt = 1e-2u"ps"
```

```
initial_system = generate_atoms_in_cubic_cell(N, element, box_size, reference_temp)
```

```
eq_steps = 2000
```





```
eq_thermostat = Molly.AndersenThermostat(reference_temp, Δt * coupling_factor)
eq_simulator = MollySimulator(Δt, eq_steps, coupling = eq_thermostat)
potential = InteratomicPotentials.LennardJones(austrip(1.657e-21u"J"),
        austrip(0.34u"nm"),
        austrip(0.765u"nm"), [:Ar])
```


```
eq_result = simulate(initial_system, eq_simulator, potential)
```

```
prod_steps = 5000
```

```
prod_simulator = MollySimulator(Δt, prod_steps, t₀ = get_time(eq_result))
```

```
prod_result = simulate(get_system(eq_result), prod_simulator, potential)
```





## Ab Initio MD

```
using Atomistic
using DFTK
using InteratomicPotentials
using NBodySimulator
```

```
setup_threading(n_blas = 4)
```

```
N = 8
element = :Ar
box_size = 1.5u"nm" # this number was chosen arbitrarily
reference_temp = 94.4u"K"
thermostat_prob = 0.1 # this number was chosen arbitrarily
Δt = 1e-2u"ps"
```

```
initial_system = generate_atoms_in_cubic_cell(N, element, box_size, reference_temp)
pspkey = list_psp(:Ar, functional = "lda")[1].identifier
for atom ∈ initial_system
    atom.data[:pseudopotential] = pspkey
end
```

```
eq_steps = 20000
eq_thermostat = NBodySimulator.AndersenThermostat(austrip(reference_temp),
                                                    thermostat_prob / austrip(Δt))
eq_simulator = NBSimulator(Δt, eq_steps, thermostat = eq_thermostat)
potential = InteratomicPotentials.LennardJones(austrip(1.657e-21u"J"),
                                                austrip(0.34u"nm"),
                                                austrip(0.765u"nm"), [:Ar])
```

```
eq_result = simulate(initial_system, eq_simulator, potential)
```

```
ab_initio_steps = 200
ab_initio_simulator = NBSimulator(Δt, ab_initio_steps, t₀ = get_time(eq_result))
dftk_potential = DFTKPotential(5u"hartree", [1, 1, 1]; damping = 0.7)
```

```
ab_initio_result = simulate(get_system(eq_result), ab_initio_simulator, dftk_potential)
```

## Atomistic.jl

```
using Atomistic
using DFTK
using InteratomicPotentials
using Molly
```

```
setup_threading(n_blas = 4)
```

```
N = 8
element = :Ar
box_size = 1.5u"nm" # this number was chosen arbitrarily
reference_temp = 94.4u"K"
coupling_factor = 10 # this number was chosen arbitrarily
Δt = 1e-2u"ps"
```

```
initial_system = generate_atoms_in_cubic_cell(N, element, box_size, reference_temp)
pspkey = list_psp(:Ar, functional = "lda")[1].identifier
for atom ∈ initial_system
    atom.data[:pseudopotential] = pspkey
end
```

```
eq_steps = 20000
eq_thermostat = Molly.AndersenThermostat(reference_temp, Δt * coupling_factor)
eq_simulator = MollySimulator(Δt, eq_steps, coupling = eq_thermostat)
potential = InteratomicPotentials.LennardJones(austrip(1.657e-21u"J"),
                                                austrip(0.34u"nm"),
                                                austrip(0.765u"nm"), [:Ar])
```

```
eq_result = simulate(initial_system, eq_simulator, potential)
```

```
ab_initio_steps = 200
ab_initio_simulator = MollySimulator(Δt, ab_initio_steps, t₀ = get_time(eq_result))
dftk_potential = DFTKPotential(5u"hartree", [1, 1, 1]; damping = 0.7)
```

```
ab_initio_result = simulate(get_system(eq_result), ab_initio_simulator, dftk_potential)
```



## ❖ Goal



- ❖ Allow **direct interaction with LAMMPS** without round-tripping through Python

## ❖ Achievements

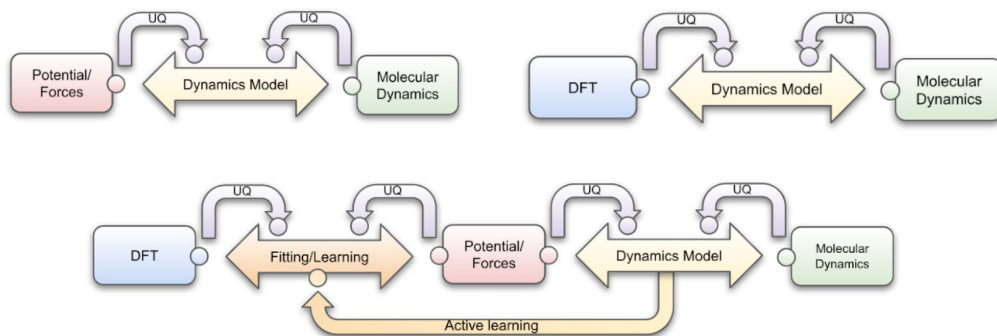
- ❖ Precompiled binary available through LAMMPS.jl
  - ❖ A binary released as a Julia package that the package manager can install
  - ❖ Transparently loading dependencies and makes libraries available
- ❖ Automatic conversion to Julia arrays (zero-cost access)
- ❖ API coverage sufficient for simple SNAP example

## ❖ Next steps

- ❖ **Integration with MPI.jl**
- ❖ **Calling Julia potential from LAMMPS**
- ❖ Increasing API coverage
- ❖ Spack integration for binary deployment on HPC systems
- ❖ Documentation and more examples
- ❖ Deeper integration than possible through the C-API



# Atomistic Composable Workflows



## ❖ General goal

- ❖ Gather easy-to-use CESMIX-aligned case studies
- ❖ Integrate the latest developments of CESMIX and the Julia atomistic ecosystem with state-of-the-art tools.

CW	Type	DFT	UQ	Fitting/ Learning	Potential / Forces	Molecular Dynamics
1	Ar		✓	✓	InteratomicPotentials.jl → Lennard Jones / ACE	LAMMPS.jl → LAMMPS
1	Ar		✓	✓	InteratomicPotentials.jl → LennardJones / ACE	Atomistic.jl → Molly.jl
1	Na				LAMMPS.jl → EAM	LAMMPS.jl → LAMMPS
3	Na	DFTK.jl	✓	✓	InteratomicPotentials.jl → SNAP / ACE	LAMMPS.jl → LAMMPS
3	HfO <sub>2</sub> , a-HfO <sub>2</sub>	QE		✓	InteratomicBasisPotentials.jl → ACE → ACE1.jl	Atomistic.jl → Molly.jl

# Other CESMIX tools and future integration efforts

## ❖ *Enzyme.jl*

- AD of statically analyzable LLVM
- Meet or exceed the performance of state-of-the-art AD tools.

## ❖ *MLP.jl: ML potentials*

- Fit and evaluate POD/SNAP/ACE potentials
- Generate and select configurations to perform DFT calculations to obtain DFT data for training potentials
- Develop NN and GNN potentials by leveraging Enzyme.jl for the differentiation of the loss function

## ❖ *MDP.jl*


- Run MD simulations for interatomic potentials that are fitted with MLP.jl
- Potentials will interact with InteratomicPotentials.jl

## ❖ *KMC.jl*

- Run Kinetic Monte Carlo in Julia.
- Will be built to interact with InteratomicPotentials.jl

## ❖ *Carom.jl*

- Interacting Particle Systems for Bayesian inference.



# Software engineering practices

## ❖ Composable Software


- ❖ *High level software that is readable, flexible, and extensible.*
- ❖ *Deep integration: generic programming, abstractions, multiple dispatch.*

## ❖ Development platform: **Github**

- ❖ **Internet *hosting***
- ❖ **Version control: *Git***
  - *Used for coordinating work among programmers*
  - *Tracks and provides control over changes to source code*
  - *Speed, data integrity, and support for distributed, non-linear workflows (i.e., thousands of parallel branches running on different systems)*
- ❖ **Continuous integration (CI)**
  - *Continuously build and test the code to make sure that the commit doesn't introduce errors*
- ❖ **Documentation**
  - *Docstring: A package for building documentation from docstrings and markdown files*
- ❖ **Open Source**


## ❖ Other tools

- ❖ *Visual Studio Code: Live Share*
- ❖ *Pkg: Julia's package manager*
- ❖ *Pluto notebooks: Julia notebooks*



# Integrating using notebooks...

- ❖ *Reactive, lightweight, and easy Julia notebooks*
- ❖ *Easy entry point*
- ❖ *Programming directly from the **browser***
- ❖ *Condenses **documentation/tutorials, code** and **live results***
  - ❖ *Can be exported as html and used as a **blogpost***
- ❖ *It can be used as a **mini-laboratory** to experiment with modifying small parts of the code*
- ❖ *This feature is experimental: you can **run** the notebook **in the cloud**, so you do not need to install the software on your computer*




# External software dependencies

## ❖ Current software dependencies

- ❖ LAMMPS: large-scale atomic/molecular massively parallel simulator
- ❖ QE: electronic-structure calculations and materials modeling
- ❖ ACE.jl: atomic cluster expansion
- ❖ Molly.jl: movement and interaction of molecules
- ❖ NBodySimulator.jl: simulating systems of N interacting bodies
- ❖ Unitful.jl: physical units
- ❖ GalacticOptim.jl: global optimization package
- ❖ Flux.jl, Lux.jl: ML frameworks
- ❖ DataLoaders.jl: threaded data iterator for ML on out-of-memory datasets
- ❖ BSON: binary JSON serialization format
- ❖ CUDA.jl: NVIDIA GPUs support
- ❖ Zygote.jl: AD
- ❖ More!

## ❖ Potential future software dependencies

- ❖ GraphNeuralNetwork.jl: matches DGL and PyTorch Geometric
- ❖ FluxGeometric.jl: GNNs
- ❖ Chemellia: ML ecosystem for systems made of atoms
- ❖ Clustering.jl: data clustering
- ❖ More!



# Thanks!... questions?

